Satwik **Bhattamishra**

PhD student, University of Oxford

🌐 satwikb.com   ✉ satwik55@gmail.com   ⌂ github.com/satwik77   🎓 Google Scholar

Education

Present Oct 2021	**University of Oxford** Ph.D. in Computer Science · Google DeepMind Scholarship Advisors: Prof. Phil Blunsom and Prof. Varun Kanade	**Oxford, UK**
May 2019 Aug 2014	**Birla Institute of Technology and Science Pilani** B.E. (Hons.), Computer Science and Int. M.Sc. (Hons.), Biological Science	**Pilani, India**

Industry Experience

Mar 2025 Dec 2024	**Google** *Student Researcher	Cloud AI Research Team* Explored scaling test-time compute and training for LLM agents. Explored various decoding and self-verification strategies to improve per-step action prediction accuracy of agents. Tools: PyTorch, GPUs	**Sunnyvale, US**
Oct 2024 June 2024	**Cohere** *Intern of Technical Staff	Foundations team* Pretrained LLMs with *hybrid* architectures involving SSMs and Transformers. Pretrained 7B models, optimized inference strategies, & integrated μP for hyperparameter transfer. Tools: JAX, TPUs	**London, UK**
Oct 2023 June 2023	*Intern of Technical Staff	Foundations team* Pretrained and evaluated LLMs with state-space and long-convolutional architectures. Pretrained 7B models and conducted ablations to evaluate the efficacy of architectural components. Tools: JAX, TPUs	
July 2021 July 2019	**Microsoft Research** *Research Fellow	Advisors: Dr. Navin Goyal, Dr. Monojit Choudhury* Analyzed the computational capabilities of Transformers and LSTMs to model various classes of formal languages; worked on compositional generalization and math word problem solving. Tools: PyTorch, GPUs	**Bangalore, India**

Publications

Provably Learning Attention with Queries   [paper]
Satwik Bhattamishra, Kulin Shah, Michael Hahn, Varun Kanade
2026 Preprint Under Review   [Preprint]

Automata Learning and Identification of the Support of Language Models   [paper]
Satwik Bhattamishra, Michael Hahn, Varun Kanade
2026 International Conference on Learning Representations   [ICLR'26]

Benefits and Limitations of Communication in Multi-Agent Reasoning   [paper]
Michael Rizvi-Martel, Satwik Bhattamishra, Neil Rathi, Guillaume Rabusseau, Michael Hahn
2026 International Conference on Learning Representations   [ICLR'26]

The Transformer Cookbook   [paper]
A. Yang, C. Watson, A. Xue, Satwik Bhattamishra, J. Llarena, W. Merrill, E. Ferreira, A. Svete, D. Chiang
2026 Transactions on Machine Learning Research   [TMLR'26]

A Formal Framework for Understanding Length Generalization in Transformers   [paper]
X. Huang, A. Yang, Satwik Bhattamishra, Y. Sarrof, A. Krebs, H. Zhou, P. Nakkiran, M. Hahn
2025 International Conference on Learning Representations   [ICLR'25]

Separations in the Representational Capabilities of Transformers and Recurrent Architectures   [paper]
Satwik Bhattamishra, Michael Hahn, Phil Blunsom, Varun Kanade
2024 Conference on Neural Information Processing Systems   [NeurIPS'24]

Understanding In-Context Learning in Transformers by Learning to Learn Discrete Functions   [paper]
Satwik Bhattamishra, Arkil Patel, Phil Blunsom, Varun Kanade
2024 International Conference on Learning Representations **[Oral, Top 1.2%]**   [ICLR'24]

Inshrinkerator: Compressing Deep Learning Training Checkpoints via Dynamic Quantization   [paper]
A. Agrawal, S. Reddy, Satwik Bhattamishra, V. Sarath Nookala, V. Vashishth, K. Rong, A. Tumanov
2024 ACM Symposium on Cloud Computing   [SoCC'24]

Simplicity Bias in Transformers and their Ability to Learn Sparse Boolean Functions [paper]
Satwik Bhattamishra, Arkil Patel, Varun Kanade, Phil Blunsom
2023 Annual Meeting of the Association for Computational Linguistics `[ACL'23]`

MAGNIFICo: Evaluating the In-Context Learning Ability of Large Language Models to Generalize to Novel Interpretations [paper]
Arkil Patel, Satwik Bhattamishra, Siva Reddy, Dzmitry Bahdanau
2023 Conference on Empirical Methods in Natural Language Processing [Oral] `[EMNLP'23]`

Revisiting the Compositional Generalization Abilities of Neural Sequence Models [paper]
Arkil Patel, Satwik Bhattamishra, Phil Blunsom, Navin Goyal
2022 Annual Meeting of the Association for Computational Linguistics `[ACL'22]`

Are NLP Models really able to Solve Simple Math Word Problems? [paper]
Arkil Patel, Satwik Bhattamishra, Navin Goyal
2021 Conference of the North American Chapter of the Association for Computational Linguistics `[NAACL'21]`

On the Ability and Limitations of Transformers to Recognize Formal Languages [paper]
Satwik Bhattamishra, Kabir Ahuja, Navin Goyal
2020 Conference on Empirical Methods in Natural Language Processing `[EMNLP'20]`

On the Practical Ability of Recurrent Neural Networks to Recognize Hierarchical Languages [paper]
Satwik Bhattamishra, Kabir Ahuja, Navin Goyal
2020 International Conference on Computational Linguistics [Best Short Paper Award] `[COLING'20]`

On the Computational Power of Transformers and Its Implication in Sequence Modeling [paper]
Satwik Bhattamishra, Arkil Patel, Navin Goyal
2020 Conference on Computational Natural Language Learning `[CoNLL'20]`

Unsung Challenges of Building Language Technologies for Low Resource Language Communities [paper]
P. Joshi, C. Barnes, S. Santy, S. Khanuja, S. Shah, A. Srinivasan, Satwik Bhattamishra, S. Sitaram, M. Choudhury, K. Bali
16th International Conference on Natural Language Processing `[ICON'19]`

Submodular Optimization-based Diverse Paraphrasing and its Effectiveness in Data Augmentation [paper]
Ashutosh Kumar*, Satwik Bhattamishra*, Manik Bhandari, Partha Talukdar (* = Equal Contribution)
2019 Conference of the North American Chapter of the Association for Computational Linguistics [Oral] `[NAACL'19]`

Selected Dev Projects

Student Developer   Google Summer of Code (GSoC) 2016   May'16 - Aug'16

Worked on the Phenopacket scraper project, which extracts information from texts scraped from life sciences websites, analyzes them, and generates a phenopacket based on the correct ontology references. [Project Page]

Machine Learning Contests [Kaggle Profile]

Kaggle Level: Competitions Expert. **Silver** medal in Kaggle Santander Value Prediction Challenge, Rank: Pvt. 185th | Pub. 189th/4484. **Bronze** medal in Kaggle Instacart Competition, Rank: Pvt. 195th | Pub. 74th/2623. Qualified for **Zonal** Round in India Hacks Machine Learning Competition by HackerEarth, Rank: 29th/860.

Review Miner | Microsoft Code.Fun.Do Hackathon [GitHub]

Developed a cross-platform application which analyses reviews from commercial websites and provides insights about products based on keyword extraction and sentiment analysis. **Winner** of Hackathon at BITS Pilani, Rank 1st/90+ teams.

Selected Talks

Representational Capabilities of Transformers and Recurrent Architectures   [Video]
  > Charles University, Czech Republic [10/25] | Saarland University, Germany [07/25]

Language Modelling with Recurrent and State Space Architectures   [Video]
  > Georgia Tech (SysML Guest Lecture) [11/24]

On the Ability of Neural Sequence Models to Recognize Formal Languages   [Slides]
  > DeepMind NLP Reading Group [03/22] | FLaNN [02/22] | MALL Lab, IISc [12/20]

Teaching and Services

Teaching Assistant ×2   **Computational Learning Theory, Oxford**   Fall 22 and Fall 23
Teaching Assistant   **Neural Networks and Fuzzy Logic, BITS Pilani**   Jan'18 - May'18
Reviewer   ICML 2024, 2023; NeurIPS 2023, 2022; ACL 2023, 2022; ICLR 2022; EMNLP 2022, 2021, 2020; NAACL 2021